

Data-Oriented Django

Adam Johnson

Computers are friggin' fast

3.2 billion operations per second

3.2 billion operations per second
× 10 cores

$$\begin{aligned} & 3.2 \text{ billion operations per second} \\ & \quad \times 10 \text{ cores} \\ & = 32 \text{ billion operations per second} \end{aligned}$$

Grug-brained Adam

1419122974736928279
+ 9350264782187121390

Grug-brained Adam

1419122974736928279
+ 9350264782187121390

55 seconds...

Grug-brained Adam

$$\begin{array}{r} 1419122974736928279 \\ + 9350264782187121390 \end{array}$$

55 seconds...

$$= 10769387766924049669$$

Grug-brained Adam

$$\begin{array}{r} 1419122974736928279 \\ + 9350264782187121390 \end{array}$$

55 seconds...

$$= 10769387766924049669$$



Adam: 0.018 ops/sec, 0% accuracy

Macbook: 32 billion ops/sec, 100% accuracy

Where does all that speed go?

Data-Oriented Design

It's not...

It's not...

- Domain-Driven Design

It's not...

- Domain-Driven Design
- Data-Driven Design

It's not...


- Domain-Driven Design
- Data-Driven Design
- Data-Oriented Programming

Data-Oriented Design



Shoshin: Beginner's mind





Software's only job is to
transform data

**Users only care about
*getting their output data***

Data Characteristics

- Format
- Volume
- Latency
- Throughput
- Statistical distribution

Context is Everything

por exemplo

Check if a number exists in a given set

por exemplo

Check if a number exists in a given set

```
numbers: set[int]  
  
def is_match(number: int) -> bool:  
    return number in numbers
```

What if there's just a single number?

What if there's just a single number?

```
def is_match(number: int) -> bool:  
    return number == 42
```

What if a billion numbers?

What if a billion numbers?

```
def is_match(number: int) -> bool:
    with connection.cursor():
        cursor.execute(
            "SELECT 1 FROM numbers WHERE n = %s",
            (number, ),
        )
        return (cursor.fetchone() is not None)
```

**What if matching many numbers against
few?**

What if matching many numbers against few?

```
numbers: set[int]

def matches(sought: set[int]) -> set[int]:
    return numbers.intersection(sought)
```

What if set is all odd numbers between 1 and 1 million?

What if false positives are acceptable? (bloom filter)


...

**Implementation depends on data
characteristics**

**Software's only job is to
transform data using
*specific hardware***


What hardware?

What hardware?




TI-85?

What hardware?



Quantum Computer?

What hardware?




Turing Tumble?

What hardware?




General-purpose CPU


General-purpose CPU




Processor-memory gap



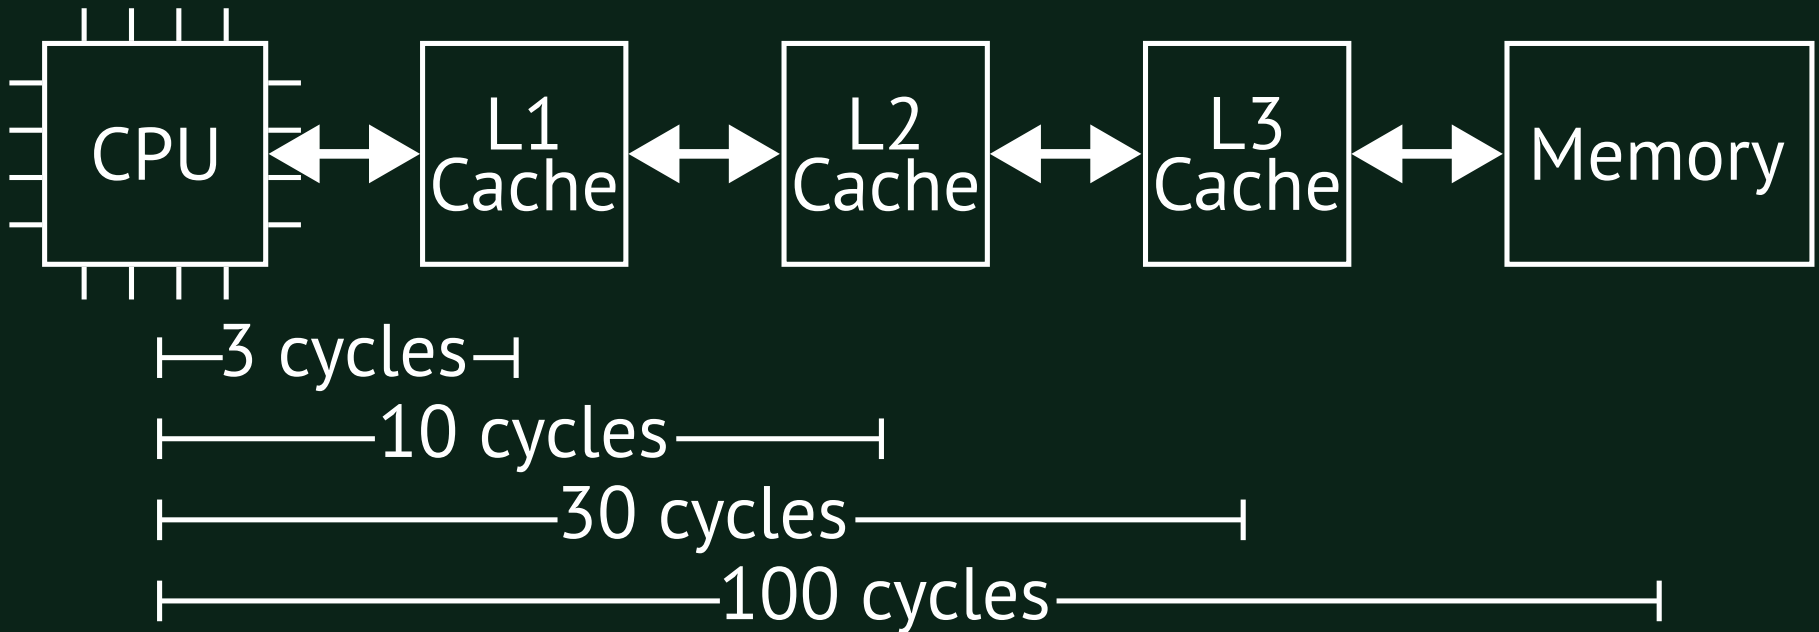
CPU caches



Three levels of cache!



Three levels of cache!



Implications

- Use smaller representations
- Lay out data in access order

Python's built-in types are big

```
In [1]: 64 // 8
```

```
Out[1]: 8
```

```
In [2]: import sys
```

```
In [3]: sys.getsizeof(9001)
```

```
Out[3]: 28
```

Python's built-in types are big

```
In [4]: numbers = list(range(1000))
```

```
In [5]: len(numbers) * 8
```

```
Out[5]: 8000
```

```
In [6]: (
...:     sys.getsizeof(numbers)
...:     + sum(sys.getsizeof(n) for n in numbers)
...: )
```

```
Out[6]: 36052
```

C data types much slimmer

```
In [7]: import array
```

```
In [8]: numbers2 = array.array('Q', range(1000))
```

```
In [9]: sys.getsizeof(numbers2)
```

```
Out[9]: 8320
```


C data types much slimmer

```
In [7]: import array
```


```
In [8]: numbers2 = array.array('Q', range(1000))
```

```
In [9]: sys.getsizeof(numbers2)
```


```
Out[9]: 8320
```

Normally use numpy or pandas for arrays

Data-Oriented Design *for the Web*

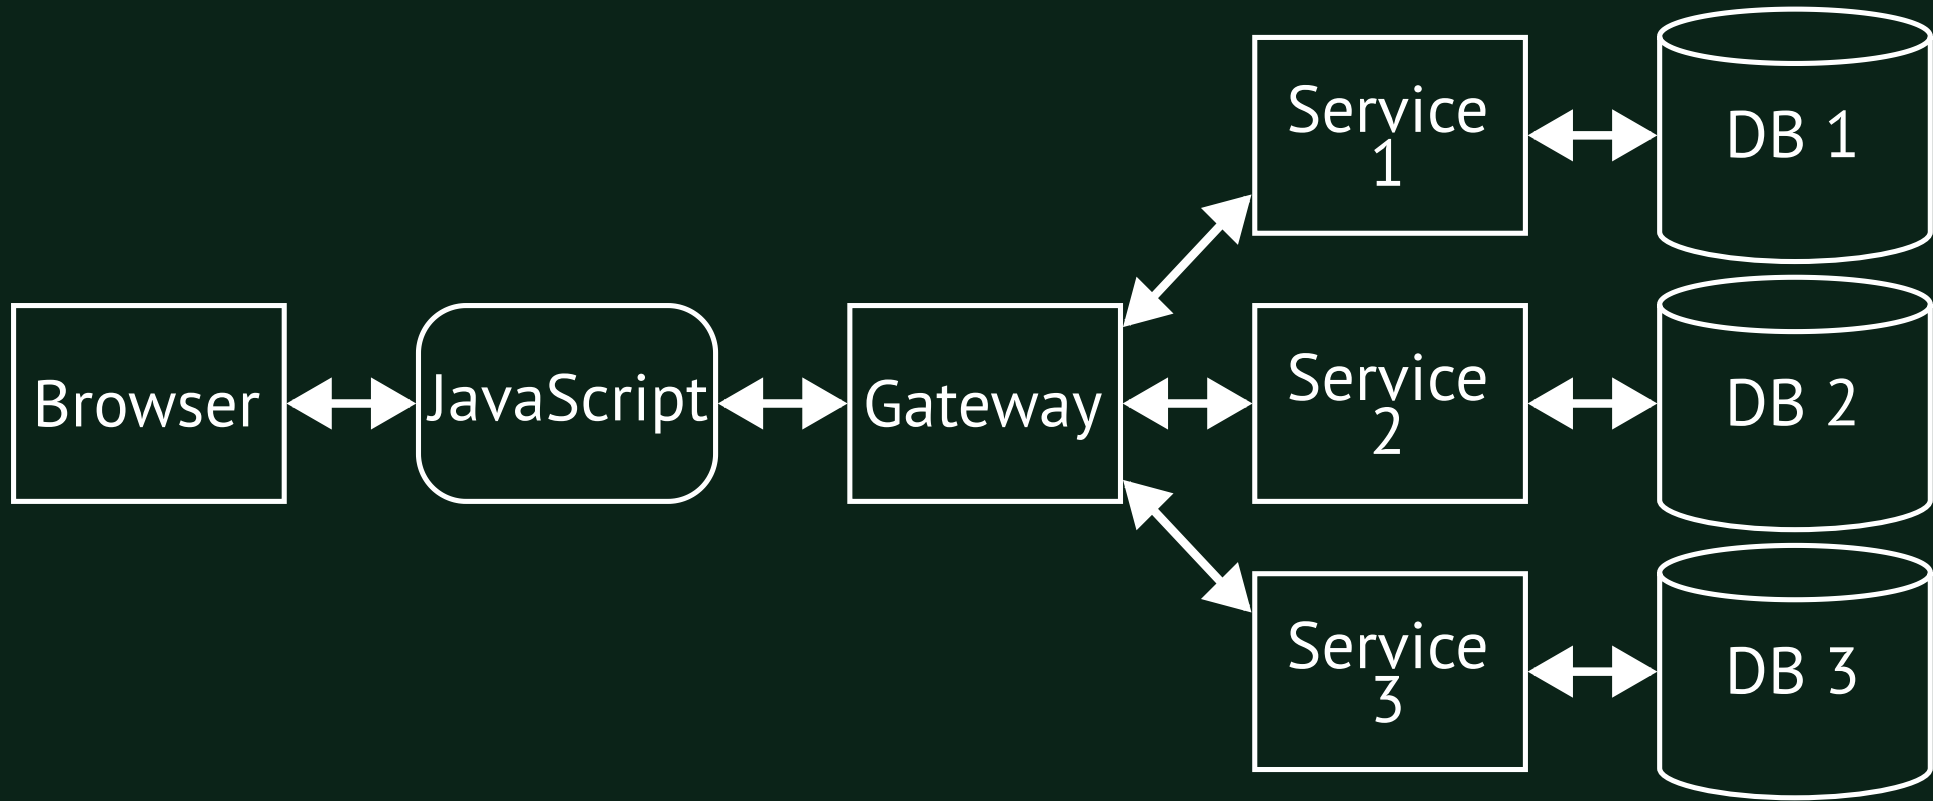


already hella fast










How fast?

$\leq 100\text{ms}$ - 🏆 🏆

$\leq 1\text{s}$ - 🏆

$\leq 3\text{s}$ - 🙌 - at 3s lose ~50% of visitors

$> 10\text{s}$ - users likely to retry, or give up



Speed up request/response cycle

Speed up request/response cycle

1. Write minimal, performant HTML

Speed up request/response cycle

1. Write minimal, performant HTML
2. HTTP Caching

Speed up request/response cycle

1. Write minimal, performant HTML
2. HTTP Caching
3. HTTP/3, or at least 2

Speed up request/response cycle

1. Write minimal, performant HTML
2. HTTP Caching
3. HTTP/3, or at least 2
4. Response compression (GZipMiddleware)

Speed up request/response cycle


1. Write minimal, performant HTML
2. HTTP Caching
3. HTTP/3, or at least 2
4. Response compression (GZipMiddleware)
5. HTML minification (django-minify-html)

Speed up request/response cycle


1. Write minimal, performant HTML
2. HTTP Caching
3. HTTP/3, or at least 2
4. Response compression (GZipMiddleware)
5. HTML minification (django-minify-html)
6. Use a CDN (Content Delivery Network)

Resources

- MDN: Web Performance
- web.dev/learn
- [WebPageTest.org](https://webpagetest.org)
- web.dev/measure



Speed up query/result cycle



Speed up queries and results

Speed up queries and results

1. Avoid N+1 queries

Speed up queries and results

1. Avoid N+1 queries
2. Split models

Speed up queries and results

1. Avoid N+1 queries
2. Split models
3. Multiple counts in one pass

1. Avoid N+1 queries

```
books = Book.objects.order_by("title")
for book in books:
    print(book.title, "by", book.author.name)
```

1. Avoid N+1 queries

```
books = Book.objects.order_by("title")
for book in books:
    print(book.title, "by", book.author.name)
```

1. Iterate books
2. For each of **N** books:
 fetch `book.author`

1. Avoid N+1 queries

```
books = (  
    Book.objects.order_by("title")  
    .select_related("author")  
)  
for book in books:  
    print(book.title, "by", book.author.name)
```

1. Avoid N+1 queries

```
books = (  
    Book.objects.order_by("title")  
    .select_related("author")  
)  
for book in books:  
    print(book.title, "by", book.author.name)
```

1. Fetch books with author joined in

1. Avoid N+1 queries

book.name	author.name
The Hundred and One Dalmatians	Dodie Smith
The Lost World	Arthur Conan Doyle
The Hound of the Baskervilles	Arthur Conan Doyle
His Last Bow	Arthur Conan Doyle

1. Avoid N+1 queries

```
books = (  
    Book.objects.order_by("title")  
    .prefetch_related("author")  
)  
for book in books:  
    print(book.title, "by", book.author.name)
```

1. Avoid N+1 queries

```
books = (  
    Book.objects.order_by("title")  
    .prefetch_related("author")  
)  
for book in books:  
    print(book.title, "by", book.author.name)
```

1. Fetch books

2. Fetch related authors for all books

1. Avoid N+1 queries

django-auto-prefetch

```
books = Book.objects.order_by("title")
for book in books:
    print(book.title, "by", book.author.name)
```


1. Avoid N+1 queries

django-auto-prefetch

```
books = Book.objects.order_by("title")
for book in books:
    print(book.title, "by", book.author.name)
```

1. Fetch books

2. On first access of `book.author`:

Fetch related authors for all books

2. Split models

```
class User(AbstractUser):  
    avatar = models.ImageField(...)  
    ...
```

2. Split models

```
class User(AbstractUser):  
    avatar = models.ImageField(...)  
    ...
```

Task: *Store user's ACME access token and refresh token*

2. Split models

Not great:

```
class User(AbstractUser):  
    avatar = models.ImageField(...)  
    ...  
    acme_access_token = models.TextField()  
    acme_access_expires = models.DateTimeField()  
    acme_refresh_token = models.TextField()
```

2. Split models

Not great:

```
class User(AbstractUser):  
    avatar = models.ImageField(...)  
    ...  
    acme_access_token = models.TextField()  
    acme_access_expires = models.DateTimeField()  
    acme_refresh_token = models.TextField()
```

Slows down every place users are queried

2. Split models

```
class User(AbstractUser):  
    avatar = models.ImageField(...)  
    ...  
  
class UserAcmeToken(models.Model):  
    user = models.OneToOneField(User, primary_key=True)  
    access_token = models.TextField()  
    access_expires = models.DateTimeField()  
    refresh_token = models.TextField()
```



3. Multiple counts in one pass

```
published_count = (  
    author.book_set.filter(verified=True).count()  
)  
unpublished_count = (  
    author.book_set.filter(verified=False).count()  
)
```

3. Multiple counts in one pass

```
counts = (  
    author.book_set.aggregate(  
        verified=Count('pk', filter=Q(verified=True)),  
        unverified=Count('pk', filter=Q(verified=False)),  
    )  
)
```


Resources

- Docs: Database access optimization
- The Temple of Django Database Performance
- Post: “Django and the N+1 Queries Problem”
- django-debug-toolbar (or Kolo)

Data-Oriented Design

Software's only job is to
transform data

**Users only care about
*getting their output data***

Resources

- Mike Acton - Data-Oriented Design and C++
- Andrew Kelley - Practical DOD
- Andreas Fredriksson - Context is Everything

Thank you! 🙌

- Adam Johnson
- @adamchainz on GitHub & Twitter
- me@adamj.eu
- github.com/adamchainz/talk-data-oriented-django
- Books: **Boost Your Django DX & Speed Up Your Django Tests**